

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score,
KFold, GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor, StackingRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MinMaxScaler

# Load datasets
fed_files = ['CPIAUCSL.csv', 'RRVRUSQ156N.csv', 'MORTGAGE30US.csv',
'MSPUS.csv',
            'USAUCSFRCONDOSMSAMID.csv', 'GDPC1.csv', 'UNRATE.csv',
            'PI.csv',
            'POPTOTUSA647NWDB.csv', 'CSUSHPIISA.csv']
data_frames = [pd.read_csv(f, parse_dates=True, index_col=0) for f in
fed_files]

data = pd.concat(data_frames, axis=1)
data = data.ffill().dropna()

data.columns = ['CPI', 'RentalVacancyRate', 'MortgageRate',
'MedianSalePrice', 'ZillowIndex',
                'GDP', 'UnemploymentRate', 'PersonalIncome',
'Population', 'HomePriceIndex']
data

```

	CPI	RentalVacancyRate	MortgageRate	MedianSalePrice
\				
DATE				
2004-01-08	186.300	10.4	5.87	212700.0
2004-01-15	186.300	10.4	5.66	212700.0
2004-01-22	186.300	10.4	5.64	212700.0
2004-01-29	186.300	10.4	5.68	212700.0
2004-02-01	186.700	10.4	5.68	212700.0
...
2024-03-21	312.230	6.6	6.87	426800.0
2024-03-28	312.230	6.6	6.79	426800.0
2024-04-01	313.207	6.6	6.79	412300.0

2024-05-01	313.207	6.6	6.79	412300.0
2024-06-01	313.207	6.6	6.79	412300.0
\	ZillowIndex	GDP	UnemploymentRate	PersonalIncome
	DATE			
2004-01-08	163387.988490	15248.680	5.7	9731.8
2004-01-15	163387.988490	15248.680	5.7	9731.8
2004-01-22	163387.988490	15248.680	5.7	9731.8
2004-01-29	163387.988490	15248.680	5.7	9731.8
2004-02-01	164448.866854	15248.680	5.6	9765.4
...
2024-03-21	359543.076444	22758.752	3.8	23746.6
2024-03-28	359543.076444	22758.752	3.8	23746.6
2024-04-01	361420.150687	22918.739	3.9	23809.6
2024-05-01	361420.150687	22918.739	4.0	23923.7
2024-06-01	361420.150687	22918.739	4.1	23923.7
DATE	Population	HomePriceIndex		
2004-01-08	292805298.0	141.647		
2004-01-15	292805298.0	141.647		
2004-01-22	292805298.0	141.647		
2004-01-29	292805298.0	141.647		
2004-02-01	292805298.0	143.192		
...		
2024-03-21	334914895.0	318.217		
2024-03-28	334914895.0	318.217		
2024-04-01	334914895.0	319.048		
2024-05-01	334914895.0	319.048		
2024-06-01	334914895.0	319.048		
[1268 rows x 10 columns]				

```

# Ensure the date index has frequency information
data.index = pd.to_datetime(data.index)
data = data.asfreq('MS') # Monthly Start frequency

# Create new feature: normalized home price index
data['NormalizedHomePriceIndex'] = data['HomePriceIndex'] /
data['CPI']

# Add lag features
for lag in range(1, 13):
    data[f'HomePriceIndex_Lag_{lag}'] =
data['HomePriceIndex'].shift(lag)

# Drop rows with missing values created by lag features
data = data.dropna()

# Verify if data is correctly processed
print(f"Data after preprocessing: {data.shape}")
print(data.head())

```

Data after preprocessing: (233, 23)

	CPI	RentalVacancyRate	MortgageRate	MedianSalePrice	\
DATE					
2005-02-01	192.4	10.1	5.66	232500.0	
2005-03-01	193.1	10.1	5.69	232500.0	
2005-04-01	193.7	9.8	6.04	233700.0	
2005-05-01	193.6	9.8	5.78	233700.0	
2005-06-01	193.7	9.8	5.65	233700.0	

	ZillowIndex	GDP	UnemploymentRate	PersonalIncome	\
DATE					
2005-02-01	182971.283783	15844.727	5.4	10282.7	
2005-03-01	184533.914455	15844.727	5.2	10347.9	
2005-04-01	186294.280193	15922.782	5.2	10412.1	
2005-05-01	188164.303058	15922.782	5.1	10464.1	
2005-06-01	190122.978622	15922.782	5.0	10498.0	

	Population	HomePriceIndex	...	HomePriceIndex_Lag_3	\
DATE			...		
2005-02-01	295516599.0	163.344	...	157.528	
2005-03-01	295516599.0	165.813	...	159.331	
2005-04-01	295516599.0	167.502	...	161.289	
2005-05-01	295516599.0	169.351	...	163.344	
2005-06-01	295516599.0	171.191	...	165.813	

	HomePriceIndex_Lag_4	HomePriceIndex_Lag_5
HomePriceIndex_Lag_6 \		
DATE		

2005-02-01	155.752	154.180
152.634		
2005-03-01	157.528	155.752
154.180		
2005-04-01	159.331	157.528
155.752		
2005-05-01	161.289	159.331
157.528		
2005-06-01	163.344	161.289
159.331		

	HomePriceIndex_Lag_7	HomePriceIndex_Lag_8
HomePriceIndex_Lag_9 \		
DATE		

2005-02-01	151.338	149.851
148.186		
2005-03-01	152.634	151.338
149.851		
2005-04-01	154.180	152.634
151.338		
2005-05-01	155.752	154.180
152.634		
2005-06-01	157.528	155.752
154.180		

	HomePriceIndex_Lag_10	HomePriceIndex_Lag_11 \
DATE		

2005-02-01	146.592	145.058
2005-03-01	148.186	146.592
2005-04-01	149.851	148.186
2005-05-01	151.338	149.851
2005-06-01	152.634	151.338

	HomePriceIndex_Lag_12
DATE	

2005-02-01	143.192
2005-03-01	145.058
2005-04-01	146.592
2005-05-01	148.186
2005-06-01	149.851

[5 rows x 23 columns]

```

# Features and target variable
X = data.drop(columns=['HomePriceIndex'])
y = data['HomePriceIndex']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Initialize models
lr = LinearRegression()
rf = RandomForestRegressor(n_estimators=100, random_state=42)
gbr = GradientBoostingRegressor(random_state=42)

# Initialize k-fold cross-validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Cross-validation for Linear Regression
cv_scores_lr = cross_val_score(lr, X, y, cv=kf,
scoring='neg_mean_squared_error')
mse_cv_lr = -cv_scores_lr.mean()

# Cross-validation for Random Forest
cv_scores_rf = cross_val_score(rf, X, y, cv=kf,
scoring='neg_mean_squared_error')
mse_cv_rf = -cv_scores_rf.mean()

# Fit models on the entire training set
lr.fit(X_train, y_train)
rf.fit(X_train, y_train)
gbr.fit(X_train, y_train)

# Predict on the test set
y_pred_lr = lr.predict(X_test)
y_pred_rf = rf.predict(X_test)
y_pred_gbr = gbr.predict(X_test)

# Ensemble predictions by averaging
y_pred_ensemble = (y_pred_lr + y_pred_rf) / 2

# Stacking Regressor
param_grid_rf = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

grid_search_rf = GridSearchCV(estimator=rf, param_grid=param_grid_rf,
cv=kf, scoring='neg_mean_squared_error', n_jobs=-1, verbose=2)
grid_search_rf.fit(X_train, y_train)

```

```

best_rf = grid_search_rf.best_estimator_

param_grid_gbr = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 5, 7]
}

grid_search_gbr = GridSearchCV(estimator=gbr,
    param_grid=param_grid_gbr, cv=kf, scoring='neg_mean_squared_error',
    n_jobs=-1, verbose=2)
grid_search_gbr.fit(X_train, y_train)
best_gbr = grid_search_gbr.best_estimator_

estimators = [
    ('lr', lr),
    ('rf', best_rf),
    ('gbr', best_gbr)
]
stacking_regressor = StackingRegressor(estimators=estimators,
    final_estimator=LinearRegression())
stacking_regressor.fit(X_train, y_train)

```

Predictions

```
y_pred_stacking = stacking_regressor.predict(X_test)
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits

Fitting 5 folds for each of 27 candidates, totalling 135 fits

Evaluation

```

def evaluate_model(y_true, y_pred, model_name):
    mse = mean_squared_error(y_true, y_pred)
    r2 = r2_score(y_true, y_pred)
    mae = mean_absolute_error(y_true, y_pred)
    print(f'{model_name} - MSE: {mse:.4f}')
    print(f'{model_name} - R2: {r2:.4f}')
    print(f'{model_name} - MAE: {mae:.4f}')

```

Evaluate each model

```

evaluate_model(y_test, y_pred_lr, 'Linear Regression')
evaluate_model(y_test, y_pred_rf, 'Random Forest')
evaluate_model(y_test, y_pred_gbr, 'Gradient Boosting')
evaluate_model(y_test, y_pred_ensemble, 'Ensemble Model')
evaluate_model(y_test, y_pred_stacking, 'Stacking Ensemble Model')

```

Linear Regression - MSE: 0.1508

Linear Regression - R2: 0.9999

Linear Regression - MAE: 0.2783

Random Forest - MSE: 1.8930


Random Forest - R2: 0.9993

```
Random Forest - MAE: 0.8921
Gradient Boosting - MSE: 2.3872
Gradient Boosting - R2: 0.9991
Gradient Boosting - MAE: 1.1100
Ensemble Model - MSE: 0.4940
Ensemble Model - R2: 0.9998
Ensemble Model - MAE: 0.4716
Stacking Ensemble Model - MSE: 0.1342
Stacking Ensemble Model - R2: 0.9999
Stacking Ensemble Model - MAE: 0.2581
```

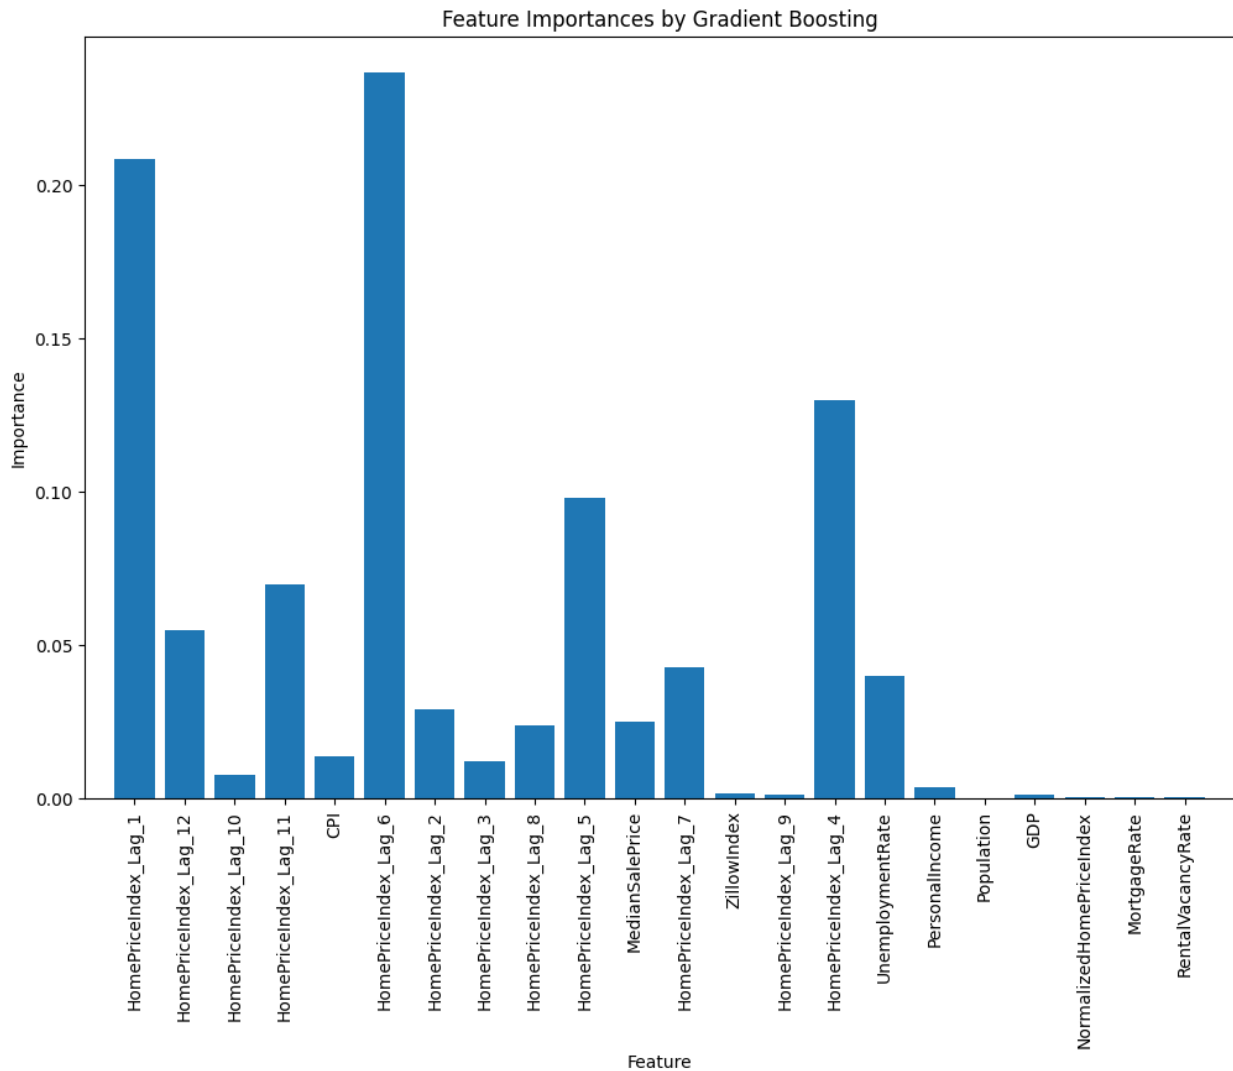
```
# Plot feature importances for Random Forest
```

```
plt.figure(figsize=(12, 8))
importances = best_rf.feature_importances_
features = X.columns
indices = np.argsort(importances)[::-1]

plt.title('Feature Importances by Random Forest')
plt.bar(range(X.shape[1]), importances[indices], align='center')
plt.xticks(range(X.shape[1]), features[indices], rotation=90)
plt.xlim([-1, X.shape[1]])
plt.xlabel('Feature')
plt.ylabel('Importance')
plt.show()
```

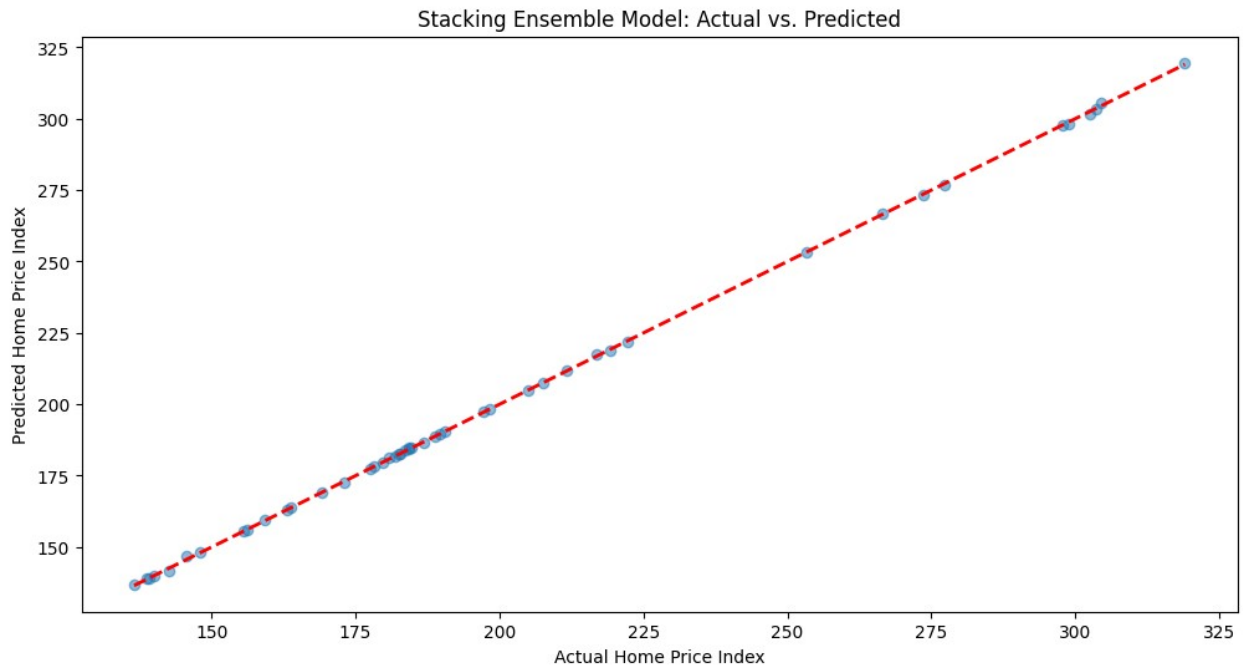


```
# Plot feature importances for Gradient Boosting
plt.figure(figsize=(12, 8))
importances = best_gbr.feature_importances_
plt.title('Feature Importances by Gradient Boosting')
plt.bar(range(X.shape[1]), importances[indices], align='center')
plt.xticks(range(X.shape[1]), features[indices], rotation=90)
plt.xlim([-1, X.shape[1]])
plt.xlabel('Feature')
plt.ylabel('Importance')
plt.show()
```

The above plot represents the feature importance of Gradient Boosting model

```
# Plot predictions vs. actual values for Stacking Ensemble Model
plt.figure(figsize=(12, 6))
plt.scatter(y_test, y_pred_stacking, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'r--', lw=2)
plt.xlabel('Actual Home Price Index')
plt.ylabel('Predicted Home Price Index')
plt.title('Stacking Ensemble Model: Actual vs. Predicted')
plt.show()
```



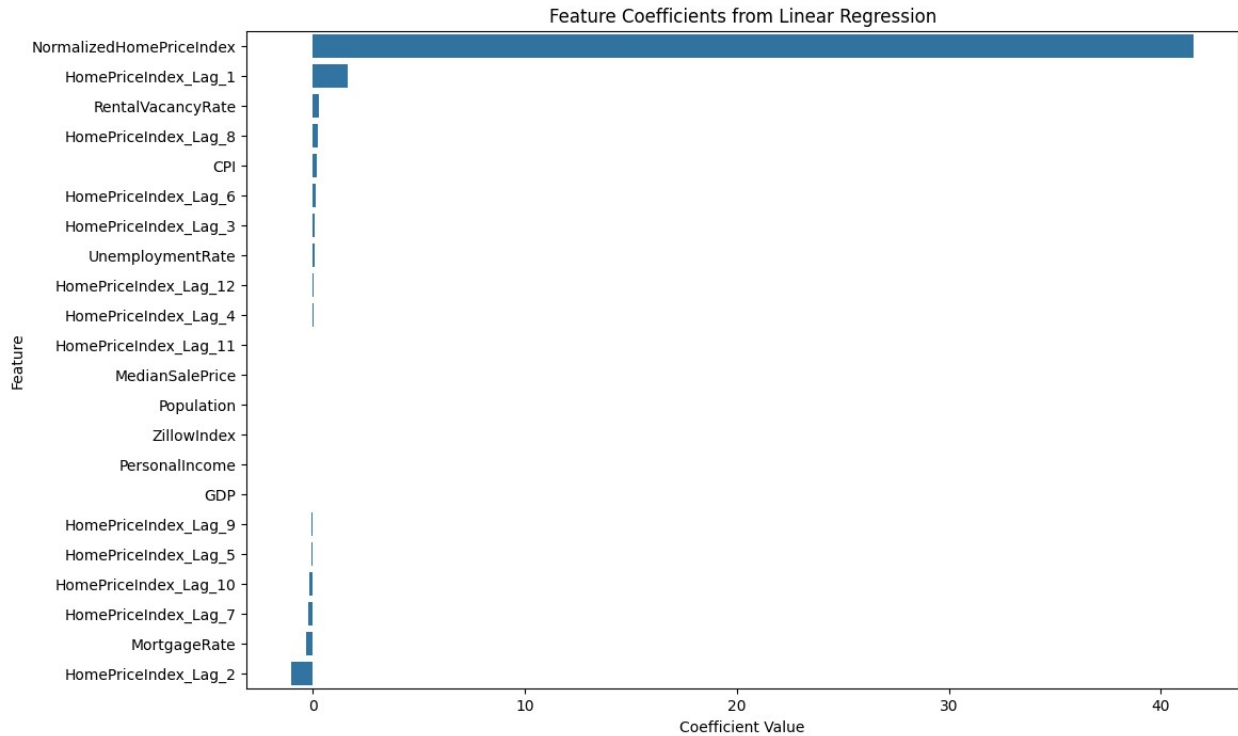
```
import matplotlib.pyplot as plt
import seaborn as sns

# Fit the Linear Regression model
lr = LinearRegression()
lr.fit(X_train, y_train)

# Get feature coefficients
coefficients = lr.coef_
features = X.columns

# Create a DataFrame for better visualization
coef_df = pd.DataFrame({
    'Feature': features,
    'Coefficient': coefficients
}).sort_values(by='Coefficient', ascending=False)


# Plot feature coefficients
plt.figure(figsize=(12, 8))
sns.barplot(x='Coefficient', y='Feature', data=coef_df)
plt.title('Feature Coefficients from Linear Regression')
plt.xlabel('Coefficient Value')
plt.ylabel('Feature')
plt.show()
```




```
# Predict on the test set
y_pred_lr = lr.predict(X_test)

# Calculate residuals
residuals = y_test - y_pred_lr


# Plot Residuals
plt.figure(figsize=(10, 6))
plt.scatter(y_pred_lr, residuals, alpha=0.5)
plt.axhline(y=0, color='r', linestyle='--')
plt.title('Residual Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.show()
```



```
# Plot Actual vs. Predicted Values
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_lr, alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
color='r', linestyle='--')
plt.title('Actual vs. Predicted Values')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.show()
```

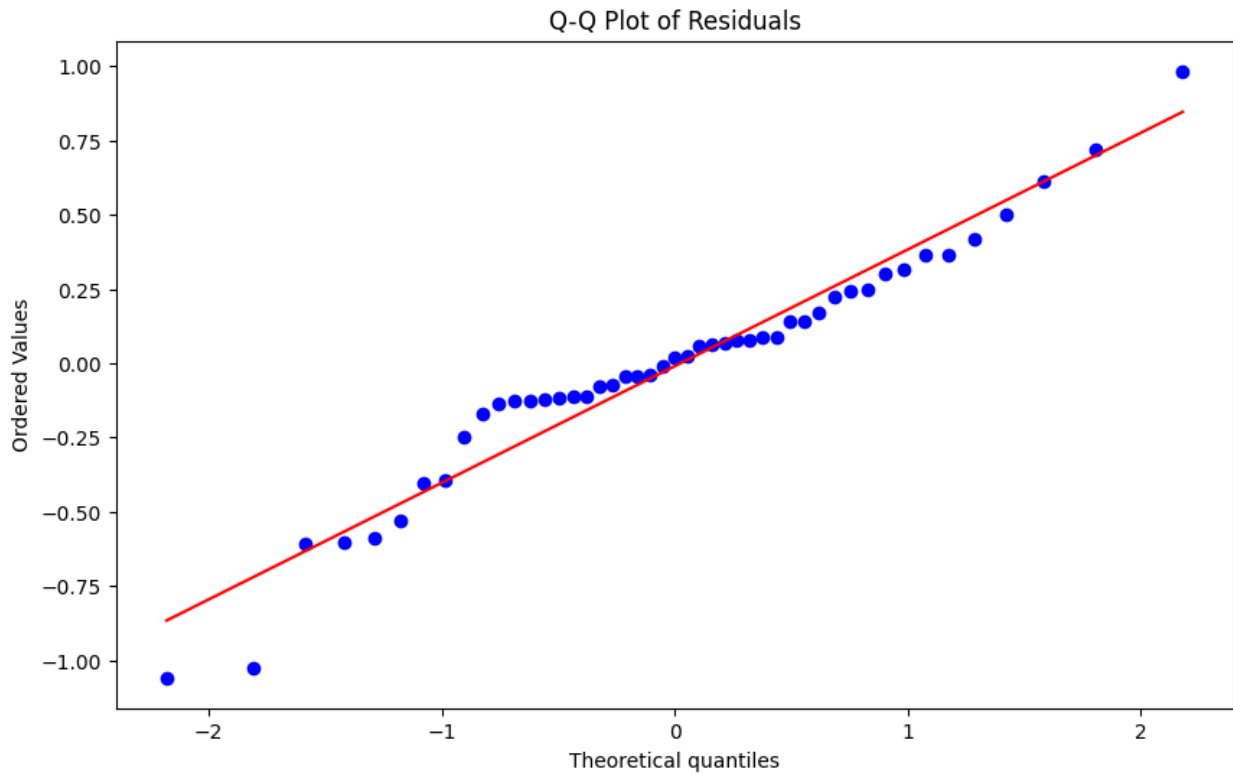


```
# Plot Distribution of Residuals
plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True, bins=30)
plt.title('Distribution of Residuals')
plt.xlabel('Residual')
plt.ylabel('Frequency')
plt.show()
```



```
import scipy.stats as stats


# Q-Q Plot
plt.figure(figsize=(10, 6))
stats.probplot(residuals, dist="norm", plot=plt)
plt.title('Q-Q Plot of Residuals')
plt.show()
```



```
from sklearn.model_selection import learning_curve

# Calculate learning curves
train_sizes, train_scores, test_scores = learning_curve(lr, X, y,
cv=5, scoring='neg_mean_squared_error')


# Plot Learning Curves
plt.figure(figsize=(10, 6))
plt.plot(train_sizes, -train_scores.mean(axis=1), label='Training
Score', color='blue')
plt.plot(train_sizes, -test_scores.mean(axis=1), label='Cross-
validation Score', color='green')
plt.title('Learning Curves')
plt.xlabel('Training Set Size')
plt.ylabel('Mean Squared Error')
plt.legend()
plt.show()
```



```
import seaborn as sns

# Calculate correlation matrix
correlation_matrix = X.copy()
correlation_matrix['HomePriceIndex'] = y
correlation_matrix = correlation_matrix.corr()

# Plot Heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix[['HomePriceIndex']].sort_values(by='HomePriceIndex', ascending=False),
            annot=True, cmap='coolwarm', center=0)
plt.title('Feature Correlation with Home Price Index')
plt.show()
```

```
# Sort features by coefficient magnitude
feature_names = X.columns
sorted_indices = np.argsort(np.abs(coefficients))
top_features = feature_names[sorted_indices][-10:] # Get top 10 features
top_coefficients = coefficients[sorted_indices][-10:]

# Plot Top Features vs Target
plt.figure(figsize=(12, 8))
for i, feature in enumerate(top_features):
    plt.scatter(X[feature], y, label=feature, alpha=0.5)
plt.xlabel('Feature Value')
plt.ylabel('Home Price Index')
plt.title('Top Features vs Home Price Index')
plt.legend()
plt.show()
```

Top Features vs Home Price Index

